

PREDICTION OF AUTOMOBILE

IMPORTING THE LIBRARIES

In [1]:



```
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
from sklearn.feature_selection import RFE
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
import warnings
warnings.filterwarnings('ignore')
```

DATA PROCESSING

In [3]:



```
c_df=pd.read_csv("C:/Users/yuvak/OneDrive/Desktop/ML PROJECTS 1&2 FROM BI/AutoData (1).csv")
```

In [4]:



```
c_df.head()
```

Out [4]:

	symboling	make	fuel type	aspiration	door number	carbony	drive wheel	engine location	wh ee lb ase	carlen ght	. . .	engine size	fuel system	bo rre atio	st r o k e	co mp res sion ratio	ho rs ep ower	P e a k r p m	ci t y m p g	hi gh way m p g	p r i c e	
	0	3	alf a-romer o giulia	gas	std	two	conver tible	rw d	front	88.6	168.8	. . .	130	mpfi	3.47	2.68	9.0	111	5000	21	27	13495.0
	1	3	alf a-romer o stelvio	gas	std	two	conver tible	rw d	front	88.6	168.8	. . .	130	mpfi	3.47	2.68	9.0	111	5000	21	27	16500.0
	2	1	alf a-romer o Quadri gli o	gas	std	two	hatchback	rw d	front	94.5	171.2	. . .	152	mpfi	2.68	3.47	9.0	154	5000	19	26	16500.0
	3	2	audi 100 ls	gas	std	four	sedan	fw d	front	99.8	176.6	. . .	109	mpfi	3.19	3.40	10.0	102	5500	24	30	13950.0
	4	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	176.6	. . .	136	mpfi	3.19	3.40	8.0	115	5500	18	22	17450.0

5 rows x 25 columns

In [5]:



c_df.tail()

Out[5]:

symboling	make	fueltype	aspiration	doornumber	carbody	drivewheel	engine	helbase	carlength	engine	fuel	bo	st	compressionratio	horsepower	peakrpm	citympg	highwaympg	price		
2000	-1	volvo145e (sw)	gas	std	four	sedan	rw	front	109.1	188.8	. . .	141	mpfi	3.78	3.15	9.5	114	5400	23	28	16845.0
2001	-1	volvo144ea	gas	turbo	four	sedan	rw	front	109.1	188.8	. . .	141	mpfi	3.78	3.15	8.7	160	5300	19	25	19045.0
2002	-1	volvo244d1	gas	std	four	sedan	rw	front	109.1	188.8	. . .	173	mpfi	3.58	2.87	8.8	134	5500	18	23	21485.0
2003	-1	volvo	diesel	turbo	four	sedan	rw	front	109.1	188.8	. . .	145	idi	3.01	3.40	23.0	106	4800	26	27	2247

symboling	make	fueltype	aspiration	doornumber	carbody	drivewheel	engine location	wheelbase	carlength	carwidth	carheight	engine size	fuelsystem	boiler ratio	stroke	compression ratio	horsepower	peakrpm	citympg	highwaympg	price
2																					0.
4																					0
6																					

2																					2
0	-1																				2
4																					6
																					2
																					5.
																					0

5 rows x 25 columns

In [7]:



c_df.shape

Out[7]:

(205, 25)

In [8]:



c_df.dtypes

Out[8]:

symboling	int64
make	object
fueltype	object
aspiration	object
doornumber	object
carbody	object
drivewheel	object
engine location	object
wheelbase	float64
carlength	float64
carwidth	float64
carheight	float64

```

curbweight          int64
enginetype          object
cylindernumber      object
enginesize          int64
fuelsystem          object
boreratio           float64
stroke             float64
compressionratio    float64
horsepower          int64
peakrpm            int64
citympg            int64
highwaympg         int64
price              float64
dtype: object

```

In [9]:



`c_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 205 entries, 0 to 204
```

```
Data columns (total 25 columns):
```

#	Column	Non-Null Count	Dtype
0	symboling	205 non-null	int64
1	make	205 non-null	object
2	fueltype	205 non-null	object
3	aspiration	205 non-null	object
4	doornumber	205 non-null	object
5	carbody	205 non-null	object
6	drivewheel	205 non-null	object
7	enginelocation	205 non-null	object
8	wheelbase	205 non-null	float64
9	carlength	205 non-null	float64
10	carwidth	205 non-null	float64
11	carheight	205 non-null	float64
12	curbweight	205 non-null	int64
13	enginetype	205 non-null	object
14	cylindernumber	205 non-null	object
15	enginesize	205 non-null	int64
16	fuelsystem	205 non-null	object
17	boreratio	205 non-null	float64
18	stroke	205 non-null	float64
19	compressionratio	205 non-null	float64
20	horsepower	205 non-null	int64
21	peakrpm	205 non-null	int64

```

22  citympg          205 non-null    int64
23  highwaympg      205 non-null    int64
24  price            205 non-null    float64
dtypes: float64(8), int64(7), object(10)
memory usage: 40.2+ KB

```

In [10]:



c_df.describe()

Out[10]:

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	engineize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	0.34146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317	3.329756	3.255415	10.142537	104.117073	512.5121951	25.219512	30.751220	13276.710571
std	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693	0.270844	0.313597	3.972040	39.544167	476.985643	6.542142	6.886443	7988.852332
min	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	2.540000	2.070000	7.000000	48.000000	415.000000	13.000000	16.000000	5118.000000
25%	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000	3.150000	3.110000	8.600000	70.000000	480.000000	19.000000	25.000000	7788.000000
50%	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	3.310000	3.290000	9.000000	95.000000	520.000000	24.000000	30.000000	10295.000000
75%	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000	3.580000	3.410000	9.400000	116.000000	550.000000	30.000000	34.000000	16503.000000

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	engine size	bore ratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	price
		0000	0000			0000	0000					0000			
max	3.0000	120.9000	208.1000	72.3000	59.8000	4066.0000	326.0000	3.9400	4.1700	23.0000	288.0000	6600.0000	49.0000	54.0000	45400.0000

DATA CLEANING

In [11]:



c_df.duplicated().sum()

Out[11]:

0

In [15]:



c_df.isnull().sum()

Out[15]:

```
symboling      0
make           0
fueltype       0
aspiration     0
doornumber     0
carbody        0
drivewheel     0
engineloation  0
wheelbase      0
carlength      0
carwidth       0
carheight      0
curbweight     0
enginetype     0
cylindernumber 0
enginesize     0
fuelsystem     0
bore ratio     0
```

```
stroke          0
compressionratio 0
horsepower      0
peakrpm         0
citympg         0
highwaympg      0
price           0
dtype: int64
```

In [16]:



```
c_df['symboling'].value_counts()
```

Out[16]:

```
0    67
1    54
2    32
3    27
-1   22
-2    3
```

```
Name: symboling, dtype: int64
```


In [25]:




```
sns.pairplot(y_vars = 'symboling', x_vars = 'price', data=c_df)
```

Out[25]:

```
<seaborn.axisgrid.PairGrid at 0x11771b787f0>
```



In [26]:



```

peugeot 504      6
toyota corolla   6
toyota corona    6
subaru dl        4
mitsubishi g4    3
..
plymouth duster  1
mitsubishi mirage 1
porsche boxter   1
audi fox         1
dodge coronet custom 1
Name: make, Length: 147, dtype: int64

```

In [28]:



	sy m bol ing	m ake	f uel ty pe	as pi ra ti on	do or nu m ber	carb ody	driv e wheel	engin eloc ation	whe el base	charl en gth	.	fu elsy stem	bo rer ratio	s t ro ke	comp res sion ratio	hors epow er	peak rpm	cit ym pg	high way mpg	pri ce	car_ com pany
0	3	alf arom ero guli a	gas	std	two	con ver tible	rwd	front	88. 6	168. 8	mpfi	347	268	9.0	111	5000	21	27	13495. 0	alf arome ro

symboling	make	fuel_tpye	aspiration	doornumbr	carbony	drive_wheel	engine_location	weightlb	carlength	. . .	fuel_system	boileratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	price	car_company	
1	3	alfaromerostelvio	gas	std	two	convertible	rw	front	88.6	168.8	. . .	mpfi	3.47	2.68	9.0	111	5000	21	27	16500.0	alfaromeo
2	1	alfaromeroQuadrioglio	gas	std	two	hatchback	rw	front	94.5	171.2	. . .	mpfi	2.68	3.47	9.0	154	5000	19	26	16500.0	alfaromeo
3	2	audi100ls	gas	std	four	sedan	fw	front	99.8	176.6	. . .	mpfi	3.19	3.40	10.0	102	5500	24	30	13500.0	audi
4	2	audi100ls	gas	std	four	sedan	4wd	front	99.4	176.6	. . .	mpfi	3.19	3.40	8.0	115	5500	18	22	17500.0	audi

5 rows × 26 columns

In [29]:



```
c_df = c_df.drop(['make'], axis = 1)
```

In [30]:



```
c_df['car_company'].value_counts()
```

Out[30]:

```
toyota      31
nissan       17
mazda        15
mitsubishi  13
honda        13
subaru       12
peugeot      11
volvo         11
volkswagen   9
dodge         9
buick         8
bmw           8
audi          7
plymouth      7
saab          6
isuzu         4
porsche       4
alfa-romero   3
chevrolet     3
jaguar        3
maxda         2
vw            2
renault       2
Nissan         1
vokswagen     1
toyouta       1
mercury       1
porcshce      1
Name: car_company, dtype: int64
```

In [32]:



```
c_df['car_company'].replace('toyouta', 'toyota', inplace=True)
c_df['car_company'].replace('Nissan', 'nissan', inplace=True)
c_df['car_company'].replace('maxda', 'mazda', inplace=True)
c_df['car_company'].replace('vokswagen', 'volkswagen', inplace=True)
c_df['car_company'].replace('vw', 'volkswagen', inplace=True)
c_df['car_company'].replace('porcshce', 'porsche', inplace=True)
```

In [33]:



```
c_df['car_company'].value_counts()
```

Out[33]:

```
toyota      32
nissan       18
mazda        17
honda        13
mitsubishi  13
volkswagen  12
subaru       12
peugeot      11
volvo        11
dodge         9
bmw           8
buick         8
audi          7
plymouth      7
saab          6
porsche       5
isuzu         4
chevrolet     3
alfa-romero   3
jaguar        3
renault       2
mercury       1
```

Name: car_company, dtype: int64

mistakes are re-summerized

In [34]:




```
c_df['fueltype'].value_counts()
```

Out[34]:

```
gas      185
diesel   20
Name: fueltype, dtype: int64
```

In [35]:



```
c_df['aspiration'].value_counts()
```

Out[35]:

```
std      168
turbo     37
Name: aspiration, dtype: int64
```

In [36]:



```
c_df['doornumber'].value_counts()
```

Out[36]:

```
four      115
two        90
Name: doornumber, dtype: int64
```

In [37]:



```
c_df['carbody'].value_counts()
```

Out[37]:

```
sedan      96
hatchback  70
wagon      25
hardtop     8
convertible 6
Name: carbody, dtype: int64
```

In [38]:



```
c_df['drivewheel'].value_counts()
```

Out[38]:

```
fwd      120
rwd       76
4wd        9
Name: drivewheel, dtype: int64
```

In [39]:



```
c_df['enginelocation'].value_counts()
```

Out[39]:

```
front      202
```

```
rear      3
Name: enginelocation, dtype: int64
```

In [40]:



```
c_df['wheelbase'].value_counts().head()
```


Out[40]:

```
94.5      21
93.7      20
95.7      13
96.5       8
98.4       7
Name: wheelbase, dtype: int64
```


In [41]:



```
sns.distplot(c_df['wheelbase'],color="lime")
plt.show()
```



In [42]:



```
c_df['carlength'].value_counts().head()
```


Out[42]:

```
157.3    15
188.8    11
166.3     7
171.7     7
186.7     7
Name: carlength, dtype: int64
```

In [43]:



```
sns.distplot(c_df['carlength'],color="orange")
plt.show()
```



In [44]:



```
c_df['enginetype'].value_counts()
```

Out[44]:

```
ohc      148
ohcf      15
ohcv      13
dohc      12
l         12
rotor      4
dohcv      1
Name: enginetype, dtype: int64
```

In [45]:



```
c_df['cylindernumber'].value_counts()
```

Out[45]:

```
four      159
six        24
five       11
eight       5
two         4
three       1
twelve     1
Name: cylindernumber, dtype: int64
```

In [46]:



```
def number(x):
    return x.map({'four':4,'six':6,'five':5,'eight':8,'two':2,'three':3,'twelve':12})
c_df['cylindernumber']=c_df[['cylindernumber']].apply(number)
```

In [47]:



```
c_df['cylindernumber'].value_counts()
```

Out[47]:

```
4      159
6       24
5       11
8        5
2         4
12        1
3         1
Name: cylindernumber, dtype: int64
```

In [48]:



```
c_df['fuelsystem'].value_counts()
```

Out[48]:

```
mpfi      94
```



```
2bbl    66
idi      20
1bbl    11
spdi     9
4bbl     3
mfi      1
spfi     1
Name: fuelsystem, dtype: int64
```

Data Visulaization:

In [52]:



```
c_num = c_df.select_dtypes(include=['int64','float64'])
```

In [53]:



```
c_num.head()
```

Out [53]:

	symbo ling	wheel base	car len gth	car wi dth	car hei ght	cur bwe ight	cylind ernu mber	eng ine size	bor era tio	st ro ke	compr ession ratio	hors epo wer	pe ak rpm	cit y m pg	high way mpg	pr ice
0	3	88.6	168.8	64.1	48.8	2548	4	130	3.47	2.68	9.0	111	5000	21	27	13495.0
1	3	88.6	168.8	64.1	48.8	2548	4	130	3.47	2.68	9.0	111	5000	21	27	16500.0
2	1	94.5	171.2	65.5	52.4	2823	6	152	2.68	3.47	9.0	154	5000	19	26	16500.0
3	2	99.8	176.6	66.2	54.3	2337	4	109	3.19	3.40	10.0	102	5500	24	30	13950.0

	sy mb olin g	wh eel base	car len gth	car wi dth	car hei ght	cur bwei ght	cylind ernu mber	eng ine size	bor era tio	st roke	compr ession ratio	hors epo wer	pe ak rpm	cit y m pg	high way mpg	pr ice
4	2	99. 4	176 .6	66. 4	54. 3	282 4	5	136	3.1 9	3. 40	8.0	115	55 00	18	22	17 45 0. 0

In [54]:

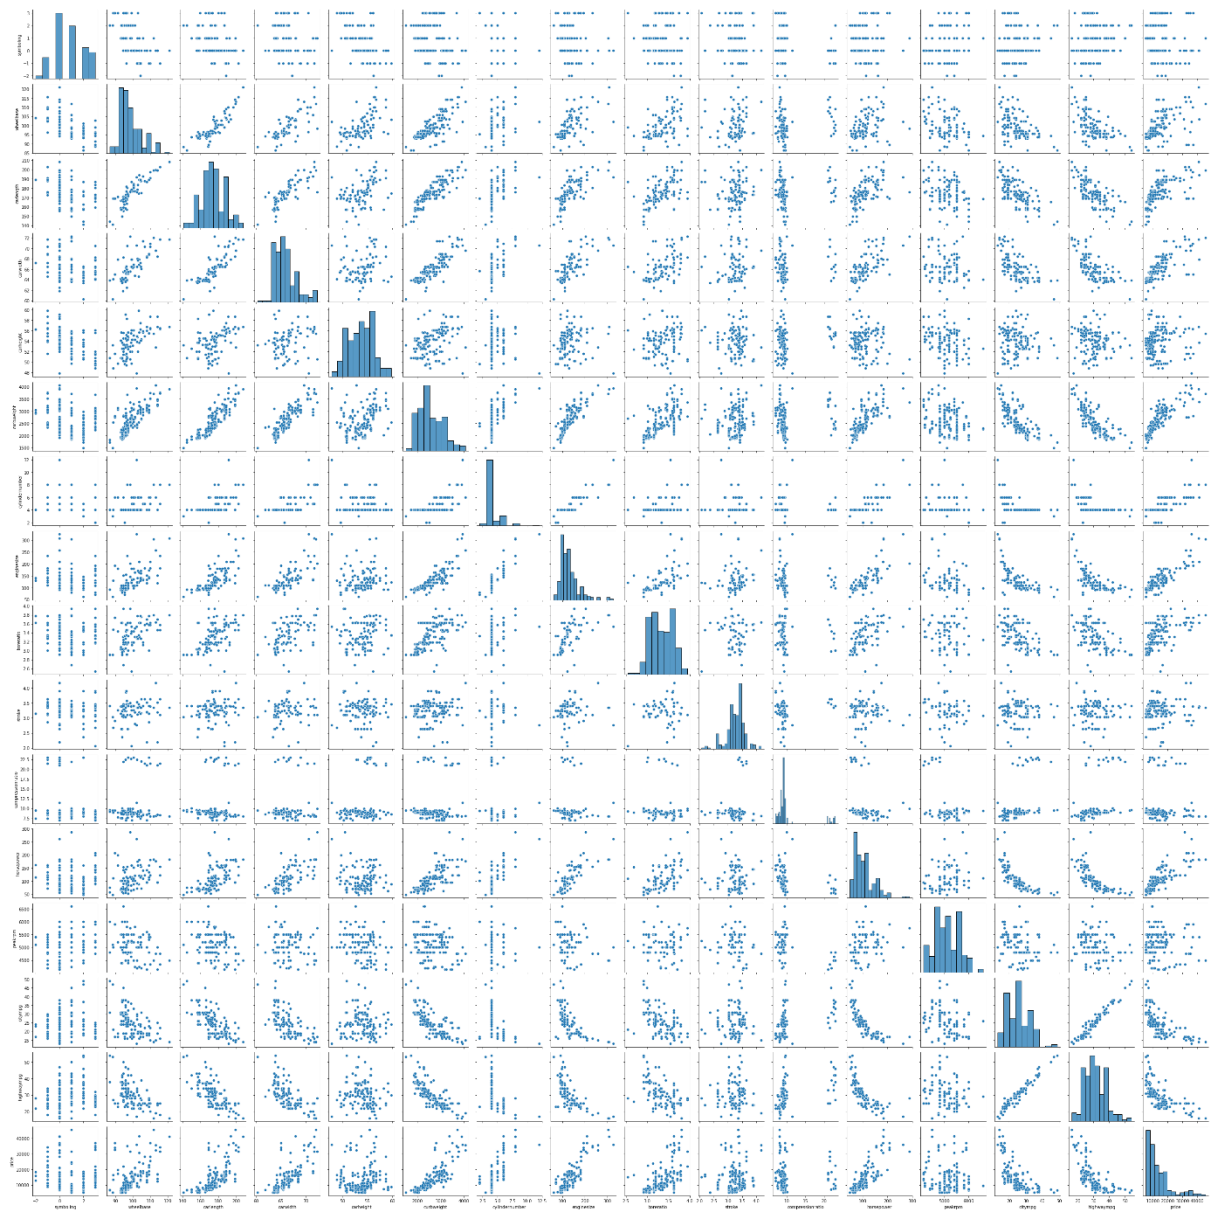


```
c_num.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   wheelbase              205 non-null    float64
2   carlength              205 non-null    float64
3   carwidth               205 non-null    float64
4   carheight              205 non-null    float64
5   curbweight             205 non-null    int64
6   cylindernumber         205 non-null    int64
7   enginesize             205 non-null    int64
8   boreratio              205 non-null    float64
9   stroke                 205 non-null    float64
10  compressionratio       205 non-null    float64
11  horsepower              205 non-null    int64
12  peakrpm                205 non-null    int64
13  citympg                205 non-null    int64
14  highwaympg             205 non-null    int64
15  price                  205 non-null    float64
dtypes: float64(8), int64(8)
memory usage: 25.8 KB
```

In [57]:




```
plt.figure(figsize = (25,25))
sns.pairplot(c_num)
plt.show()
<Figure size 1800x1800 with 0 Axes>
```



In [61]:



```
plt.figure(figsize = (25,25))
sns.heatmap(c_df.corr(), annot = True ,cmap = 'YlOrBr')
plt.show()
```




-Price is highly (positively) correlated with wheelbase, carlength, carwidth, curbweight, enginesize, horsepower.

-Price is negatively correlated to symboling, citympg and highwaympg.

-This suggest that cars having high mileage may fall in the 'economy' cars category, and are priced lower.

-There are many independent variables which are highly correlated: wheelbase, carlength, curbweight, enginesize etc.. all are positively correlated.

In [62]:



```
categorical_cols = c_df.select_dtypes(include = ['object'])
```

```
categorical_cols.head()
```

Out [62] :

	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	enginetype	fuelsystem	car_company
0	gas	std	two	convertible	rwd	front	dohc	mpfi	alfa-romero
1	gas	std	two	convertible	rwd	front	dohc	mpfi	alfa-romero
2	gas	std	two	hatchback	rwd	front	ohcv	mpfi	alfa-romero
3	gas	std	four	sedan	fwd	front	ohc	mpfi	audi
4	gas	std	four	sedan	4wd	front	ohc	mpfi	audi

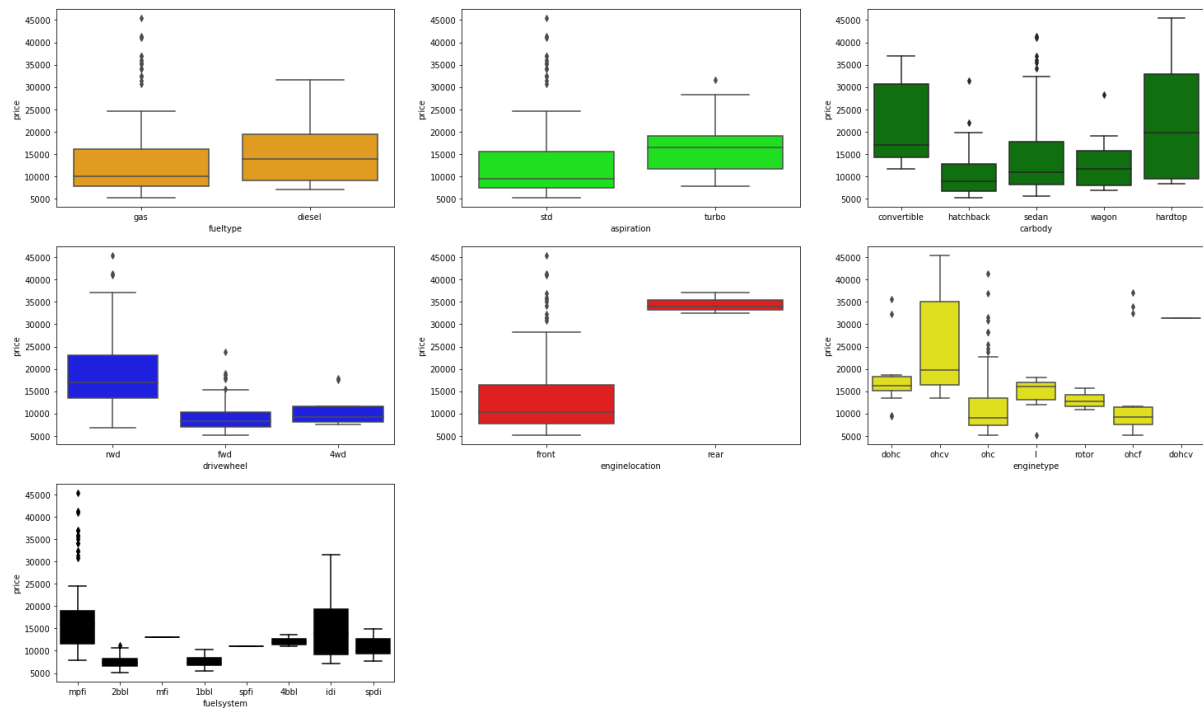
In [64] :




```
plt.figure(figsize = (25,15))
plt.subplot(3,3,1)
sns.boxplot(x = 'fueltype', y = 'price', data = c_df,color="orange")
plt.subplot(3,3,2)
sns.boxplot(x = 'aspiration', y = 'price', data = c_df,color="lime")
plt.subplot(3,3,3)
sns.boxplot(x = 'carbody', y = 'price', data = c_df,color="green")
plt.subplot(3,3,4)
sns.boxplot(x = 'drivewheel', y = 'price', data = c_df,color="blue")
plt.subplot(3,3,5)
sns.boxplot(x = 'enginelocation', y = 'price', data = c_df,color="red")
plt.subplot(3,3,6)
sns.boxplot(x = 'enginetype', y = 'price', data = c_df,color="yellow")
plt.subplot(3,3,7)
sns.boxplot(x = 'fuelsystem', y = 'price', data = c_df,color="black")
```

Out [64] :

```
<AxesSubplot:xlabel='fuelsystem', ylabel='price'>
```




In [71]:



```
plt.figure(figsize = (25,15))
sns.barplot(x = 'car_company', y = 'price', data = c_df,color="green")
```

Out[71]:

```
<AxesSubplot:xlabel='car_company', ylabel='price'>
```



1. From the price boxplot it is clear that The brands with the most expensive vehicles in the dataset belong to Bmw,Buick,Jaguar and porsche.

- ## Extract Features and Target



Out[461]:

[illegible]

s y m b o l i n g	d o o r n u m b e r	w h e e l b a s e	c a r l e n g t h	c a r w i d t h	c a r h e i g h t	c u r b w e i g h t	cy li n d e r n u m b e r	e n g i n e s i z e	b o r e a t i o n	ca r_ co m p a n y_ n i s s a n	car _c om p a n y_ p e u g e ot	car _co mp a n y_ p ly mo uth	car _c om p a n y_ p ors ch e	ca r_ co m p a n y_ r en aul t	ca r_ co m p a n y_ s aa b	ca r_ co m p a n y_ s ub ar u	ca r_ co m p a n y_ t oy ot a	car _co mp any_ vo lks wa gen	ca r_ co m p a n y_ v ol vo
		7 7	4 0	3 3		0 8		8 7	1 4										

2			0	0	0		0	0											
0	-	f										
2	1	o	6	7	7	5	5	4	7										
2		u	5	1	1	5	9	2	4	0	0	0	0	0	0	0	0	0	1
		r	5	1	6	.	1	2	2	.									
			9	9	6	5	1	6	8	.									
			7	4	6		5	4	5										
			7	0	7		6	2	7										

2			0	0	0		0	0											
0	-	f										
2	1	o	6	7	7	5	6	3	3										
0		u	5	1	1	5	7	1	3	0	0	0	0	0	0	0	0	0	1
3		r	5	1	6	.	0	6	5	.									
			9	9	6	5	6	9	7	.									
			7	4	6		7	8	1										
			7	0	7		5	1	4										

2			0	0	0		0	0											
0	-	f										
2	1	o	6	7	7	5	6	3	8										
0		u	5	1	1	5	1	0	8	0	0	0	0	0	0	0	0	0	1
4		r	5	1	6	.	0	1	5	.									
			9	9	6	5	5	8	7	.									
			7	4	6		5	8	1										
			7	0	7		1	7	4										

205 rows x 61 columns

In [462]:



y=c_df["price"]

In [463]:



X=c_df.drop("price",axis="columns")

In [464]:



X

Out[464]:

[illegible]

[illegible]

Features should be of numeric nature:

The image shows a presentation slide with a large, empty rectangular area in the center. On the right side, there are three small, square navigation buttons stacked vertically. On the bottom edge, there are two small, square navigation buttons, one on the left and one on the right. The entire slide is enclosed in a thin black border.

Out[465]:

```
NonNumericColumns=X.columns[X.dtypes=="object"]
```

In [467]:




NonNumericColumns

Out[467]:

Index(['doornumber'], dtype='object')

In [468]:



X=pd.get_dummies(X,columns=NonNumericColumns,drop_first=True)

Features should be of type array/ dataframe:

In [469]:



type(X)

Out[469]:

pandas.core.frame.DataFrame

Features should have some rows and some columns:

In [470]:



X.shape

Out[470]:

(205, 60)


Split the dataset- training and testing:

In [471]:




from sklearn.model_selection import train_test_split

In [472]:



In [473]:



Company Information										Product Details										Summary	
Symbol	Wheelbase	Car Length	Car Width	Car Height	Curb Weight	Cylinder Number	Engine Size	Bore	Stroke	Company Name	Company Address	Company Phone	Company Email	Company Website	Product Name	Product Description	Price	Availability			
ABC	100	180	70	60	3000	4	2.0	3.5	4.0	ABC Corp	123 Main St	555-1234	info@abc.com	www.abc.com	Model X	High performance sedan	\$25000	In Stock			
DEF	110	190	75	65	3200	6	2.5	3.8	4.2	DEF Inc	456 Elm St	555-5678	sales@def.com	www.def.com	Model Y	Mid-range sedan	\$18000	Low Stock			
GHI	120	200	80	70	3500	8	3.0	4.0	4.5	GHI Ltd	789 Oak St	555-9012	contact@ghi.com	www.ghi.com	Model Z	Full-size sedan	\$30000	Out of Stock			
JKL	130	210	85	75	3800	10	3.5	4.2	4.8	JKL Co	101 Pine St	555-3456	support@jkl.com	www.jkl.com	Model A	Executive sedan	\$40000	Pre-order			
MNO	140	220	90	80	4000	12	4.0	4.5	5.0	MNO Corp	202 Birch St	555-7890	info@mno.com	www.mno.com	Model B	Luxury sedan	\$50000	On Order			

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	cylindernumber	engine size	bore ratio	stroke	car_company_eugetot	car_company_plymouth	car_company_porsche	car_company_renault	car_company_saab	car_company_subaru	car_companytoyota	car_companyvolswagen	car_companyvolvo	door number_twe
min	-2	0	0	0	47	0	2.	0	0	2										
	.0	.0	.0	.0	.8	.0	.0	.0	.0	.7	0.0	0.0	0.0	0.0	0.	0.0	0.0	0.0	0.	0.
	0	0	0	0	0	0	0	0	0	0	00	00	00	00	00	00	00	000	00	00
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	00	0	0	00	00	00
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	00	0	0	00	00	00
25%	0	0	0	0	52	0	4.	0	0	3										
	.0	.2	.3	.3	.0	.2	.0	.1	.4	.1	0.0	0.0	0.0	0.0	0.	0.0	0.0	0.0	0.	0.
	0	3	7	1	0	5	0	3	3	1	00	00	00	00	00	00	00	000	00	00
	0	0	6	6	0	4	0	5	5	0	00	00	00	00	00	00	00	00	00	00
	0	3	1	6	0	8	0	8	7	0	0	0	0	0	00	0	0	00	00	00
50%	0	2	1	6	0	4	0	4	1	0										
	0	1	9	7	0	9	0	9	4	0										
	1	0	0	0	54	0	4.	0	0	3										
	.0	.3	.4	.4	.1	.3	.0	.2	.5	.2	0.0	0.0	0.0	0.0	0.	0.0	0.0	0.0	0.	0.
	0	0	7	3	0	5	0	2	5	9	00	00	00	00	00	00	00	000	00	00
75%	0	0	3	3	0	9	0	0	0	0										
	0	2	1	3	0	1	0	6	0	0	0	0	0	0	00	0	0	00	00	00
	0	0	0	3	0	9	0	4	0	0										
	0	7	4	3	0	3	0	2	0	0										
	2	0	0	0	55	0	4.	0	0	3										
max	.0	.4	.6	.5	.5	.5	.0	.3	.7	.4	0.0	0.0	0.0	0.0	0.	0.0	0.0	0.0	0.	1.
	0	6	2	5	0	6	0	0	4	1	00	00	00	00	00	00	00	000	00	00
	0	0	6	0	0	1	0	1	2	0	00	00	00	00	00	00	00	00	00	00
	0	6	8	0	0	2	0	8	8	0	0	0	0	0	00	0	0	00	00	00
	0	4	6	0	0	8	0	8	5	0	0	0	0	0	00	0	0	00	00	00

s	w	c	c	c	c	cy	e	b	s	.	ca	car	ca	ca	ca	ca	ca	ca	do
y	h	a	a	a	u	li	n	o	r	.	r	_co	_co	_co	_co	_co	_co	_co	or
m	e	r	r	r	r	n	i	r	e	.	co	_co	co	co	co	co	co	com	n
b	l	l	w	h	b	d	n	e	r	.	m	mp	mp	mp	mp	mp	mp	com	u
o	a	n	i	e	w	e	u	s	a	.	pa	an	pa	pa	pa	pa	pa	pa	be
l	s	g	d	g	i	n	m	i	t	.	_p	y_p	_p	_r	_t	_t	_t	_t	r
n	e	t	t	h	h	b	b	z	o	.	eu	mo	ors	en	oy	wa	vo	vo	tw
g						er		e		.	ge	uth	ch	aul	a				o
											ot		e						
0	0	0	0	0	0	0	0	0	0										
0	0	0	0	0	0	0	0	0	0										

8 rows x 60 columns

In [474]:




```
from sklearn.preprocessing import MinMaxScaler
```

In [475]:



```
scaler=MinMaxScaler()
```

In [476]:



```
X_train=scaler.fit_transform(X_train)
```

In [477]:



```
X_test=scaler.transform(X_test)
```

Train the model on the training dataset:

In [478]:



```
X.shape
```

Out [478]:

(205, 60)

In [479]:



```
from sklearn.feature_selection import RFE
```

In [480]:



```
from sklearn.linear_model import LinearRegression
```

In [481]:



```
model=LinearRegression()
```

In [482]:



```
rfe_model=RFE(model,15)
```

In [483]:



```
rfe_model.fit(X_train,y_train)
```

Out[483]:

```
RFE(estimator=LinearRegression(), n_features_to_select=15)
```

In [484]:



```
rfe_model.support_
```

Out[484]:

```
array([False, False, False,  True, False,  True,  True,  True,  True,
        True, False,  True, False, False, False, False, False, False,
        False, False, False, False, False, False,  True,  True, False,
        False, False, False, False, False, False, False, False, False,
        False, False,  True,  True,  True, False, False, False, False,
```

```
True, False, False, False, False, False, False, True, False,
False, False, False, False, True, False])
```

In [485]:



X.columns

Out[485]:

```
Index(['symboling', 'wheelbase', 'carlength', 'carwidth', 'carheight',
      'curbweight', 'cylindernumber', 'enginesize', 'boreratio', 'stroke',
      'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
      ,
      'fueltype_gas', 'aspiration_turbo', 'doornumber_two', 'carbody_hardt
op',
      'carbody_hatchback', 'carbody_sedan', 'carbody_wagon', 'drivewheel_f
wd',
      'drivewheel_rwd', 'engine_location_rear', 'enginetype_dohcv',
      'enginetype_l', 'enginetype_ohc', 'enginetype_ohcf', 'enginetype_ohc
v',
      'enginetype_rotor', 'fuelsystem_2bbl', 'fuelsystem_4bbl',
      'fuelsystem_idi', 'fuelsystem_mfi', 'fuelsystem_mphi',
      'fuelsystem_spdi', 'fuelsystem_spfi', 'car_company_audi',
      'car_company_bmw', 'car_company_buick', 'car_company_chevrolet',
      'car_company_dodge', 'car_company_honda', 'car_company_isuzu',
      'car_company_jaguar', 'car_company_mazda', 'car_company_mercury',
      'car_company_mitsubishi', 'car_company_nissan', 'car_company_peugeot',
      ,
      'car_company_plymouth', 'car_company_porsche', 'car_company_renault',
      ,
      'car_company_saab', 'car_company_subaru', 'car_company_toyota',
      'car_company_volkswagen', 'car_company_volvo', 'doornumber_two'],
      dtype='object')
```

In [486]:



Top15Columns=X.columns[rfe_model.support_]

In [487]:



Top15Columns

Out[487]:

```
Index(['carwidth', 'curbweight', 'cylindernumber', 'enginesize', 'boreratio',
      'stroke', 'horsepower', 'enginelocation_rear', 'enginetype_dohcv',
      'car_company_audi', 'car_company_bmw', 'car_company_buick',
      'car_company_jaguar', 'car_company_porsche', 'car_company_volvo'],
      dtype='object')
```

In [488]:




```
X_train=pd.DataFrame(X_train,columns=X.columns)
```

In [489]:



```
X_test=pd.DataFrame(X_test,columns=X.columns)
```

In [490]:



```
X_train=X_train[Top15Columns]
```

In [491]:



```
X_test=X_test[Top15Columns]
```

In [492]:



```
import statsmodels.api as sm
```

In [493]:



```
X_train
```

Out[493]:

	car width	curb weight	cylinder number	engine size	bore ratio	stroke	horse power	engine location_rear	engine type_ohev	car_ company_audi	car_ company_bmw	car_ company_buick	car_ company_jaguar	car_ company_porsche	car_ company_volvo
0	0.342857	0.1835	0.2	0.148438	0.746032	0.227273	0.1279	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.190476	0.11729	0.2	0.078125	0.230159	0.525253	0.06779	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.247619	0.098173	0.2	0.085937	0.230159	0.525253	0.06779	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.352381	0.234169	0.2	0.105469	0.261905	0.6111	0.000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.190476	0.087123	0.2	0.078125	0.230159	0.525253	0.06779	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
.
159	1.0000	0.7025	0.6	0.519531	1.0000	0.4646	1.000	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0
160	0.514286	0.564811	0.2	0.277344	0.873016	0.4848	0.62712	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
161	0.352381	0.210795	0.2	0.152344	0.404762	0.6111	0.39831	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

	car width h	cu rb we igh t	cyli nde rnu mbe r	en gi ne siz e	bo re ra tio	st ro ke	ho rse po we r	engin eloca tion_ rear	engi nety pe_d ohcv	car_ com pany_ audi	car_c omp any_ bmw	car_c ompa ny_b uick	car_c ompa ny_ja guar	car_c ompa ny_po rsche	car_c ompa ny_v olvo
1	0.	0.3		0.	0.	0.	0.0								
6	44	25	0.2	15	46	58	88	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	76	96		62	82	58	98								
	19	7		50	54	59	3								

1	0.	0.2		0.	0.	0.	0.0								
6	19	03	0.2	10	23	52	67	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	04	14		93	01	52	79								
	76	5		75	59	53	7								


164 rows x 15 columns

In [494]:



```
y_train=pd.DataFrame(y_train).reset_index(drop=True)
```

In [495]:



y_train

Out [495]:

price	
0	0.049849
1	0.061839
2	0.026588
3	0.071421
4	0.011271
...	...

price

159 0.652463

160 0.282930

161 0.070925

162 0.138523

163 0.061839

164 rows × 1 columns

In [496]:



```
pvalues=sm.OLS(y_train,X_train).fit()
```

In [497]:



```
pvalues.summary()
```

Out [497]:

OLS Regression Results

Dep. Variable:	price	R-squared (uncentered):	0.971
----------------	-------	-------------------------	-------

Model:	OLS	Adj. R-squared (uncentered):	0.968
--------	-----	------------------------------	-------

Method:	Least Squares	F-statistic:	327.3
---------	---------------	--------------	-------

Date:	Thu, 27 Jan 2022	Prob (F-statistic):	8.07e-106
-------	------------------	---------------------	-----------

Time:	09:26:34	Log-Likelihood:	261.59
-------	----------	-----------------	--------

No. Observations:	164	AIC:	-493.2
-------------------	-----	------	--------

Df Residuals:	149	BIC:	-446.7
---------------	-----	------	--------

Df Model:	15
-----------	----

Covariance Type:	nonrobust
------------------	-----------

coef	std err	t	P> t	[0.025	0.975]
------	---------	---	------	--------	--------

carwidth	0.2120	0.048	4.381	0.000	0.116	0.308
curbweight	0.1875	0.052	3.596	0.000	0.084	0.290
cylindernumber	-0.1643	0.061	-2.694	0.008	-0.285	-0.044
enginesize	0.3026	0.081	3.742	0.000	0.143	0.462
boreratio	-0.0840	0.024	-3.437	0.001	-0.132	-0.036
stroke	-0.0392	0.020	-1.926	0.056	-0.079	0.001
horsepower	0.2397	0.053	4.546	0.000	0.135	0.344
enginelocation_rear	0.2773	0.063	4.390	0.000	0.153	0.402
enginetype_dohcv	-0.0377	0.079	-0.476	0.635	-0.194	0.119
car_company_audi	0.0607	0.025	2.474	0.014	0.012	0.109
car_company_bmw	0.2757	0.024	11.414	0.000	0.228	0.323
car_company_buick	0.2533	0.032	7.948	0.000	0.190	0.316
car_company_jaguar	0.1792	0.048	3.734	0.000	0.084	0.274
car_company_porsche	0.1503	0.054	2.795	0.006	0.044	0.257
car_company_volvo	0.0635	0.022	2.914	0.004	0.020	0.107
Omnibus:	38.901	Durbin-Watson:	2.128			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	100.672			
Skew:	0.967	Prob(JB):	1.38e-22			
Kurtosis:	6.315	Cond. No.	27.8			

Notes:


- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [498]:



```
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

In [499]:




```
list(range(0,len(X_train.columns)))
```

Out[499]:

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14]
```

In [500]:



```
VIF=[]  
for i in range(len(X_train.columns)):  
    vifvalue=variance_inflation_factor(X_train.values,i)  
    VIF.append(vifvalue)  
results=pd.DataFrame()  
results["Features"]=X_train.columns  
results["VIF"]=VIF  
results.sort_values("VIF",ascending=False)
```

Out[500]:

	Features	VIF
3	enginesize	30.339968
1	curbweight	28.574975
0	carwidth	27.252337
2	cylindernumber	15.862157
6	horsepower	13.584559
4	boreratio	11.514555

	Features	VIF
5	stroke	7.957269
13	car_company_porsche	5.449883
7	enginelocation_rear	4.512491
8	enginetype_dohcv	2.362695
11	car_company_buick	2.297328
12	car_company_jaguar	1.735591
9	car_company_audi	1.363560
10	car_company_bmw	1.319310
14	car_company_volvo	1.254465

In [501]:



results

Out[501]:

	Features	VIF
0	carwidth	27.252337
1	curbweight	28.574975
2	cylindernumber	15.862157
3	enginesize	30.339968
4	boreratio	11.514555
5	stroke	7.957269

	Features	VIF
6	horsepower	13.584559
7	enginelocation_rear	4.512491
8	enginetype_dohcv	2.362695
9	car_company_audi	1.363560
10	car_company_bmw	1.319310
11	car_company_buick	2.297328
12	car_company_jaguar	1.735591
13	car_company_porsche	5.449883
14	car_company_volvo	1.254465

In [502]:



```
X_train=X_train.drop("boreratio",axis="columns")
```

In [503]:



```
X_test=X_test.drop("boreratio",axis="columns")
```

In [504]:



```
X_train.shape
```

Out [504]:

```
(164, 14)
```

In [505]:




X_test.shape

Out[505]:

(41, 14)

In [506]:



```
pvalues=sm.OLS(y_train,X_train).fit()  
pvalues.summary()
```

Out[506]:


OLS Regression Results							
Dep. Variable:		price		R-squared (uncentered):		0.968	
Model:		OLS		Adj. R-squared (uncentered):		0.965	
Method:		Least Squares		F-statistic:		326.3	
Date:		Thu, 27 Jan 2022		Prob (F-statistic):		1.27e-104	
Time:		09:28:30		Log-Likelihood:		255.34	
No. Observations:		164		AIC:		-482.7	
Df Residuals:		150		BIC:		-439.3	
Df Model:		14					
Covariance Type:		nonrobust					
		coef	std err	t	P> t	[0.025	0.975]
	carwidth	0.1668	0.048	3.460	0.001	0.072	0.262
	curbweight	0.1558	0.053	2.933	0.004	0.051	0.261
	cylindernumber	-0.1610	0.063	-2.551	0.012	-0.286	-0.036
	enginesize	0.2718	0.083	3.266	0.001	0.107	0.436
	stroke	-0.0546	0.021	-2.662	0.009	-0.095	-0.014
	horsepower	0.2345	0.055	4.297	0.000	0.127	0.342

enginelocation_rear	0.2806	0.065	4.291	0.000	0.151	0.410
enginetype_dohcv	-0.0044	0.081	-0.055	0.957	-0.165	0.156
car_company_audi	0.0877	0.024	3.642	0.000	0.040	0.135
car_company_bmw	0.2814	0.025	11.282	0.000	0.232	0.331
car_company_buick	0.2868	0.031	9.131	0.000	0.225	0.349
car_company_jaguar	0.2235	0.048	4.671	0.000	0.129	0.318
car_company_porsche	0.1269	0.055	2.298	0.023	0.018	0.236
car_company_volvo	0.0643	0.023	2.848	0.005	0.020	0.109
Omnibus:	36.709	Durbin-Watson:	2.045			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	93.492			
Skew:	0.916	Prob(JB):	4.99e-21			
Kurtosis:	6.213	Cond. No.	23.9			

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [508]:



```
VIF=[]
for i in range(len(X_train.columns)):
    vifvalue=variance_inflation_factor(X_train.values,i)
    VIF.append(vifvalue)
results=pd.DataFrame()
results["Features"]=X_train.columns
results["VIF"]=VIF
results.sort_values("VIF",ascending=False)
```

Out [508]:

	Features	VIF
3	enginesize	29.967218
1	curbweight	27.683483
0	carwidth	25.240891
2	cylindernumber	15.858262
5	horsepower	13.573386
4	stroke	7.566429
12	car_company_porsche	5.362649
6	enginelocation_rear	4.511440
7	enginetype_dohcv	2.327402
10	car_company_buick	2.082103
11	car_company_jaguar	1.610138
9	car_company_bmw	1.312935
13	car_company_volvo	1.254343
8	car_company_audi	1.224016

In [509]:



```
X_train=X_train.drop("stroke",axis="columns")
X_test=X_test.drop("stroke",axis="columns")
```

In [510]:



```
pvalues=sm.OLS(y_train,X_train).fit()  
pvalues.summary()
```

Out[510]:


OLS Regression Results							
Dep. Variable:		price		R-squared (uncentered):		0.967	
Model:		OLS		Adj. R-squared (uncentered):		0.964	
Method:		Least Squares		F-statistic:		337.2	
Date:		Thu, 27 Jan 2022		Prob (F-statistic):		2.14e-104	
Time:		09:30:59		Log-Likelihood:		251.55	
No. Observations:		164		AIC:		-477.1	
Df Residuals:		151		BIC:		-436.8	
Df Model:		13					
Covariance Type:		nonrobust					
		coef	std err	t	P> t	[0.025	0.975]
	carwidth	0.1193	0.046	2.611	0.010	0.029	0.210
	curbweight	0.1502	0.054	2.774	0.006	0.043	0.257
	cylindernumber	-0.2329	0.058	-4.003	0.000	-0.348	-0.118
	enginesize	0.2982	0.084	3.539	0.001	0.132	0.465
	horsepower	0.2396	0.056	4.307	0.000	0.130	0.350
	enginelocation_rear	0.2800	0.067	4.198	0.000	0.148	0.412
	enginetype_dohcv	0.0353	0.082	0.433	0.666	-0.126	0.196
	car_company_audi	0.0988	0.024	4.085	0.000	0.051	0.147
	car_company_bmw	0.2913	0.025	11.574	0.000	0.242	0.341

car_company_buick	0.3123	0.031	10.229	0.000	0.252	0.373
car_company_jaguar	0.2551	0.047	5.396	0.000	0.162	0.349
car_company_porsche	0.1375	0.056	2.448	0.016	0.027	0.249
car_company_volvo	0.0779	0.022	3.476	0.001	0.034	0.122
Omnibus:	35.032	Durbin-Watson:	2.024			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	90.028			
Skew:	0.869	Prob(JB):	2.82e-20			
Kurtosis:	6.187	Cond. No.	19.6			

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [511]:



```
VIF=[]
for i in range(len(X_train.columns)):
    vifvalue=variance_inflation_factor(X_train.values,i)
    VIF.append(vifvalue)
results=pd.DataFrame()
results["Features"]=X_train.columns
results["VIF"]=VIF
results.sort_values("VIF",ascending=False)
```

Out[511]:

	Features	VIF
3	enginesize	29.540020
1	curbweight	27.639910
0	carwidth	21.779889


	Features	VIF
4	horsepower	13.556504
2	cylindernumber	12.953959
11	car_company_porsche	5.334717
5	enginelocation_rear	4.511385
6	enginetype_dohcv	2.248999
9	car_company_buick	1.889715
10	car_company_jaguar	1.510843
8	car_company_bmw	1.284272
12	car_company_volvo	1.189612
7	car_company_audi	1.187264

In [513]:



```
X_train=X_train.drop("cylindernumber",axis="columns")
X_test=X_test.drop("cylindernumber",axis="columns")
```

In [514]:



```
pvalues=sm.OLS(y_train,X_train).fit()
pvalues.summary()
```

Out[514]:

OLS Regression Results			
Dep. Variable:	price	R-squared (uncentered):	0.963
Model:	OLS	Adj. R-squared (uncentered):	0.960

Method:	Least Squares			F-statistic:		331.3	
Date:	Thu, 27 Jan 2022			Prob (F-statistic):		2.28e-102	
Time:	09:34:19			Log-Likelihood:		243.28	
No. Observations:	164			AIC:		-462.6	
Df Residuals:	152			BIC:		-425.4	
Df Model:	12						
Covariance Type:	nonrobust						
	coef	std err	t	P> t	[0.025	0.975]	
carwidth	0.0519	0.045	1.166	0.246	-0.036	0.140	
curbweight	0.1954	0.056	3.520	0.001	0.086	0.305	
enginesize	0.0990	0.071	1.389	0.167	-0.042	0.240	
horsepower	0.2431	0.058	4.169	0.000	0.128	0.358	
enginelocation_rear	0.2449	0.069	3.534	0.001	0.108	0.382	
enginetype_dohcv	-0.0051	0.085	-0.060	0.952	-0.173	0.163	
car_company_audi	0.0991	0.025	3.909	0.000	0.049	0.149	
car_company_bmw	0.2945	0.026	11.170	0.000	0.242	0.347	
car_company_buick	0.3535	0.030	11.735	0.000	0.294	0.413	
car_company_jaguar	0.2716	0.049	5.501	0.000	0.174	0.369	
car_company_porsche	0.1739	0.058	2.993	0.003	0.059	0.289	
car_company_volvo	0.0869	0.023	3.716	0.000	0.041	0.133	
Omnibus:	55.065	Durbin-Watson:		2.054			
Prob(Omnibus):	0.000	Jarque-Bera (JB):		172.858			

Skew:	1.308	Prob(JB):	2.91e-38
--------------	-------	------------------	----------

Kurtosis: 7.295 **Cond. No.** 17.6

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [515]:



```
VIF=[]
for i in range(len(X_train.columns)):
    vifvalue=variance_inflation_factor(X_train.values,i)
    VIF.append(vifvalue)
results=pd.DataFrame()
results["Features"]=X_train.columns
results["VIF"]=VIF
results.sort_values("VIF",ascending=False)
```

Out [515]:

	Features	VIF
1	curbweight	26.439464
2	enginesize	19.238065
0	carwidth	18.820871
3	horsepower	13.553243
10	car_company_porsche	5.195056
4	enginelocation_rear	4.433400
5	enginetype_dohcv	2.214503
8	car_company_buick	1.674605
9	car_company_jaguar	1.499397

	Features	VIF
7	car_company_bmw	1.282924
6	car_company_audi	1.187254
11	car_company_volvo	1.177738

In [518]:



```
X_train=X_train.drop("engine size",axis="columns")
X_test=X_test.drop("engine size",axis="columns")
```

In [519]:



```
pvalues=sm.OLS(y_train,X_train).fit()
pvalues.summary()
```

Out [519]:

OLS Regression Results							
Dep. Variable:		price		R-squared (uncentered):		0.963	
Model:		OLS		Adj. R-squared (uncentered):		0.960	
Method:		Least Squares		F-statistic:		359.0	
Date:		Thu, 27 Jan 2022		Prob (F-statistic):		3.05e-103	
Time:		09:37:11		Log-Likelihood:		242.25	
No. Observations:		164		AIC:		-462.5	
Df Residuals:		153		BIC:		-428.4	
Df Model:		11					
Covariance Type:		nonrobust					
		coef	std err	t	P> t	[0.025	0.975]
	carwidth	0.0634	0.044	1.446	0.150	-0.023	0.150
	curbweight	0.2165	0.054	4.043	0.000	0.111	0.322

horsepower	0.2808	0.052	5.427	0.000	0.179	0.383
enginelocation_rear	0.2550	0.069	3.688	0.000	0.118	0.392
enginetype_dohcv	-0.0179	0.085	-0.211	0.833	-0.185	0.149
car_company_audi	0.0942	0.025	3.739	0.000	0.044	0.144
car_company_bmw	0.3037	0.026	11.863	0.000	0.253	0.354
car_company_buick	0.3727	0.027	13.877	0.000	0.320	0.426
car_company_jaguar	0.3011	0.045	6.732	0.000	0.213	0.389
car_company_porsche	0.1740	0.058	2.985	0.003	0.059	0.289
car_company_volvo	0.0829	0.023	3.560	0.000	0.037	0.129
Omnibus:	59.438	Durbin-Watson:	2.091			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	225.933			
Skew:	1.339	Prob(JB):	8.70e-50			
Kurtosis:	8.088	Cond. No.	16.3			

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [520]:




```
VIF=[]
for i in range(len(X_train.columns)):
    vifvalue=variance_inflation_factor(X_train.values,i)
    VIF.append(vifvalue)
results=pd.DataFrame()
results["Features"]=X_train.columns
results["VIF"]=VIF
```

```
results.sort_values("VIF",ascending=False)
```

Out [520] :

	Features	VIF
1	curbweight	24.457594
0	carwidth	18.166648
2	horsepower	10.610975
9	car_company_porsche	5.195050
3	enginelocation_rear	4.384744
4	enginetype_dohcv	2.188629
7	car_company_buick	1.323009
8	car_company_jaguar	1.222734
6	car_company_bmw	1.202203
5	car_company_audi	1.163552
10	car_company_volvo	1.159505

In [521] :



```
X_train=X_train.drop("curbweight",axis="columns")
X_test=X_test.drop("curbweight",axis="columns")
```

In [522] :



```
pvalues=sm.OLS(y_train,X_train).fit()
pvalues.summary()
```

Out [522] :

OLS Regression Results

Dep. Variable:	price		R-squared (uncentered):		0.959		
Model:	OLS		Adj. R-squared (uncentered):		0.956		
Method:	Least Squares		F-statistic:		357.7		
Date:	Thu, 27 Jan 2022		Prob (F-statistic):		3.63e-101		
Time:	09:38:37		Log-Likelihood:		233.93		
No. Observations:	164		AIC:		-447.9		
Df Residuals:	154		BIC:		-416.9		
Df Model:	10						
Covariance Type:	nonrobust						
	coef	std err	t	P> t	[0.025	0.975]	
carwidth	0.2034	0.028	7.197	0.000	0.148	0.259	
horsepower	0.3807	0.048	7.985	0.000	0.287	0.475	
enginelocation_rear	0.2712	0.072	3.747	0.000	0.128	0.414	
enginetype_dohcv	-0.0785	0.087	-0.899	0.370	-0.251	0.094	
car_company_audi	0.0774	0.026	2.971	0.003	0.026	0.129	
car_company_bmw	0.3161	0.027	11.860	0.000	0.263	0.369	
car_company_buick	0.3929	0.028	14.201	0.000	0.338	0.448	
car_company_jaguar	0.3309	0.046	7.155	0.000	0.240	0.422	
car_company_porsche	0.1469	0.061	2.418	0.017	0.027	0.267	
car_company_volvo	0.0902	0.024	3.708	0.000	0.042	0.138	
Omnibus:	52.980	Durbin-Watson:		2.236			
Prob(Omnibus):	0.000	Jarque-Bera (JB):		154.272			
Skew:	1.287	Prob(JB):		3.16e-34			

Kurtosis: 6.994

Cond. No. 12.9

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [524]:



```
VIF=[]
```

```
for i in range(len(X_train.columns)):
```

```
    vifvalue=variance_inflation_factor(X_train.values,i)
```

```
    VIF.append(vifvalue)
```

```
results=pd.DataFrame()
```

```
results["Features"]=X_train.columns
```

```
results["VIF"]=VIF
```

```
results.sort_values("VIF",ascending=False)
```

Out [524]:


	Features	VIF
1	horsepower	8.191734
0	carwidth	6.854279
8	car_company_porsche	5.126040
2	enginelocation_rear	4.369984
3	enginetype_dohcv	2.119965
6	car_company_buick	1.277027
7	car_company_jaguar	1.189413
5	car_company_bmw	1.184970
9	car_company_volvo	1.152389
4	car_company_audi	1.132036

In [552]:



```
X_train=X_train.drop("carwidth",axis="columns")
X_test=X_test.drop("carwidth",axis="columns")
```

In [553]:



```
pvalues=sm.OLS(y_train,X_train).fit()
pvalues.summary()
```

Out [553]:

OLS Regression Results

Dep. Variable:	price	R-squared (uncentered):	0.936			
Model:	OLS	Adj. R-squared (uncentered):	0.933			
Method:	Least Squares	F-statistic:	325.6			
Date:	Thu, 27 Jan 2022	Prob (F-statistic):	4.90e-90			
Time:	10:05:33	Log-Likelihood:	197.40			
No. Observations:	164	AIC:	-380.8			
Df Residuals:	157	BIC:	-359.1			
Df Model:	7					
Covariance Type:	nonrobust					
	coef	std err	t P> t [0.025 0.975]			
horsepower	0.7376	0.027	27.243	0.000	0.684	0.791
enginelocation_rear	0.1105	0.086	1.285	0.201	-0.059	0.280
enginetype_dohcv	-0.2203	0.106	-2.073	0.040	-0.430	-0.010
car_company_bmw	0.2653	0.032	8.226	0.000	0.202	0.329
car_company_buick	0.4265	0.032	13.219	0.000	0.363	0.490
car_company_jaguar	0.2391	0.056	4.280	0.000	0.129	0.349

car_company_porsche 0.1351 0.075 1.803 0.073 -0.013 0.283

Omnibus: 32.059 **Durbin-Watson:** 2.198

Prob(Omnibus): 0.000 **Jarque-Bera (JB):** 49.111

Skew: 1.031 **Prob(JB):** 2.17e-11

Kurtosis: 4.714 **Cond. No.** 7.53

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [554]:



```
VIF=[]
```

```
for i in range(len(X_train.columns)):
```

```
    vifvalue=variance_inflation_factor(X_train.values,i)
```

```
    VIF.append(vifvalue)
```

```
results=pd.DataFrame()
```

```
results["Features"]=X_train.columns
```

```
results["VIF"]=VIF
```


```
results.sort_values("VIF",ascending=False)
```

Out[554]:

	Features	VIF
6	car_company_porsche	5.098936
1	enginelocation_rear	4.029362
2	enginetype_dohcv	2.050239
0	horsepower	1.724928
4	car_company_buick	1.133970
5	car_company_jaguar	1.133280

	Features	VIF
3	car_company_bmw	1.132587

In [557]:




```
VIF=[]
for i in range(len(X_train.columns)):
    vifvalue=variance_inflation_factor(X_train.values,i)
    VIF.append(vifvalue)
results=pd.DataFrame()
results["Features"]=X_train.columns
results["VIF"]=VIF
results.sort_values("VIF",ascending=False)
```

Out [557]:

	Features	VIF
6	car_company_porsche	5.098936
1	enginelocation_rear	4.029362
2	enginetype_dohcv	2.050239
0	horsepower	1.724928
4	car_company_buick	1.133970
5	car_company_jaguar	1.133280
3	car_company_bmw	1.132587

In [558]:



X_train.columns

Out [558]:

```
Index(['horsepower', 'enginelocation_rear', 'enginetype_dohcv',
      'car_company_bmw', 'car_company_buick', 'car_company_jaguar',
```


symbo ling	doornu mbe r	whe el base	car len gth	car wid th	car he igh t	curb we igh t	cylin der nu m be r	eng ine dis pl a ce	bo r er a t io	ca r_ co m pa ny_ n iss an	car_ c om pa ny_ p eu ge ot	car_ co mp an y_ ply mo uth	car_ c om pa ny_ p ors ch e	ca r_ co m pa ny_ r en aul t	ca r_ co m pa ny_ s aa b	ca r_ co m pa ny_ s ub ar u	ca r_ co m pa ny_ t oy ot a	car_ co mp an y_ vo lks wa gen	ca r_ co m pa ny_ v ol vo
---------------	--------------------	-------------------	-------------------	------------------	-----------------------	------------------------	------------------------------------	------------------------------------	-------------------------------	--	---	--	---	---	--	--	--	--	---

			0	0	0		0		0	0									
											
4	2	f	3	5	5	5	5	2	4										
		o	7	2	0	4	1	8	6	.	0	0	0	0	0	0	0	0	
		u	3	9	8	.	8	3	4	.									
		r	1	8	3	3	2	0	2										
			7	5	3		3	1	8										
			8	1	3		1	9	6										


5 rows x 61 columns

In [560]:



model=LinearRegression()

In [561]:




model.fit(X_train,y_train)

Out[561]:

LinearRegression()

In [562]:




model.coef_

Out[562]:

```
array([[ 0.68554384,  0.12465035, -0.18828237,  0.27183568,  0.4332186 ,
         0.2613242 ,  0.14057525]])
```

In [563]:



```
results=pd.DataFrame()
results["Features"]=X_train.columns
```

In [564]:



X_train.columns

Out[564]:

```
Index(['horsepower', 'enginelocation_rear', 'enginetype_dohcv',
      'car_company_bmw', 'car_company_buick', 'car_company_jaguar',
      'car_company_porsche'],
      dtype='object')
```

In [565]:




X_train.columns

Out[565]:

```
Index(['horsepower', 'enginelocation_rear', 'enginetype_dohcv',
      'car_company_bmw', 'car_company_buick', 'car_company_jaguar',
      'car_company_porsche'],
      dtype='object')
```

In [568]:



```
results["Importance"]=model.coef_.reshape(7,)
```

In [569]:



```
results.sort_values(by="Importance",ascending=False)
```

Out[569]:

	Features	Importance
0	horsepower	0.685544
4	car_company_buick	0.433219

	Features	Importance
3	car_company_bmw	0.271836
5	car_company_jaguar	0.261324
6	car_company_porsche	0.140575
1	enginelocation_rear	0.124650
2	enginetype_dohcv	-0.188282

Here, above are the features that makes car price huge among all car drives.